

CASE STUDY- INTRODUCTION TO KUBERNETES

You have just joined a startup Ventura Software as a DevOps Lead Engineer. The company relies on a Monolithic Architecture for its product. Recently, the senior management was hired. The new CTO insists on having a Microservice Architecture. The Development Team is working on breaking the Monolith. Meanwhile, you have been asked to host a Test Application on Kubernetes, to understand how it works.

The following things must be implemented:

1. Deploy an Apache2 deployment of 2 replicas
2. Sample code has been checked in at the following Git-Hub repo:

<https://github.com/hshar/website.git>.

You must containerize this code and push it to Docker Hub. Once done, deploy it on Kubernetes with 2 replicas

3. Deploy Ingress with the following rules:

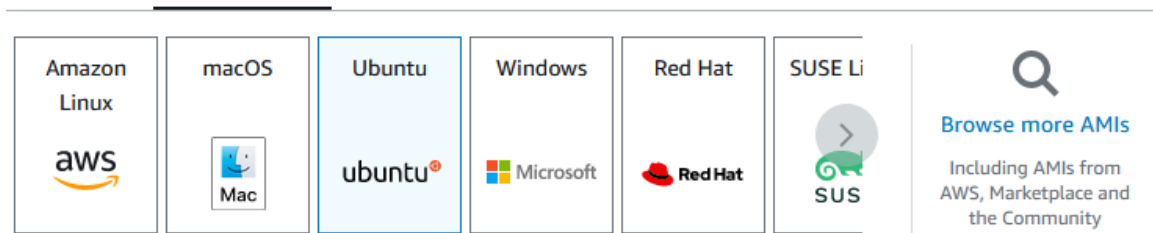
i) `*/Apache*` should point to the Apache pods

ii) `*/custom*` should point to the GitHub application

Answer:

As I am using AWS resources to create a Kubernetes environment, I am using AWS EC2 instances Ubuntu 20.04 LTS, launching t2.medium.

Login to your AWS console and navigate to EC2 service → Click on Launch Instance



Amazon Machine Image (AMI)

Ubuntu Server 20.04 LTS (HVM), SSD Volume Type

Free tier eligible ▼

ami-0a7cf821b91bccbc (64-bit (x86)) / ami-025a235c91853ccbe (64-bit (Arm))

Virtualization: hvm ENA enabled: true Root device type: ebs

Description

Canonical, Ubuntu, 20.04 LTS, amd64 focal image build on 2023-10-25

Architecture

64-bit (x86) ▼

AMI ID

ami-0a7cf821b91bccbc

Verified provider

▼ Instance type [Info](#) | [Get advice](#)

Instance type

t2.medium

Family: t2 2 vCPU 4 GiB Memory Current generation: true

On-Demand Linux base pricing: 0.0496 USD per Hour

On-Demand Windows base pricing: 0.0676 USD per Hour

On-Demand RHEL base pricing: 0.1096 USD per Hour

On-Demand SUSE base pricing: 0.1496 USD per Hour

▼

☐ All generations

[Compare instance types](#)

[Additional costs apply for AMIs with pre-installed software](#)

- Select OS: Ubuntu 20.04LTS Version
- Instance Type: t2.medium (Cost will incur)
- Select Keypair (If you do not have a key pair please create one)
- Network Settings Allow all the traffic to avoid issues

Description - *required* | [Info](#)

launch-wizard-2 created 2024-01-11T09:56:57.969Z

Inbound Security Group Rules

▼ Security group rule 1 (All, All, 0.0.0.0/0)

Remove

Type | [Info](#)

All traffic ▼

Protocol | [Info](#)

All

Port range | [Info](#)

All

Source type | [Info](#)

Anywhere ▼

Source | [Info](#)

🔍 Add CIDR, prefix list or security

0.0.0.0/0 ✕

Description - *optional* | [Info](#)

e.g. SSH for admin desktop

⚠️ Rules with source of 0.0.0.0/0 allow all IP addresses to access your instance. We recommend setting security group rules to allow access from known IP addresses only. ✕

Add security group rule

Example like above.

Then click on Launch Instance.

Launch instance

Wait for a few mins since it will take 3 to 5 mins to prep the instance.

<input type="checkbox"/>	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS
<input type="checkbox"/>	KubeInstance	i-06a973cc9d4f9956c	Running	t2.medium	2/2 checks passed	View alarms +	ap-south-1a	ec2-13-233-168-118.ap...

Once you see the Instance Checks are green then try to access the instance.

Now in the Linux terminal type the below commands to prep the environment.

```
$ sudo apt-get update
$ sudo apt-get install docker.io -y
$ curl -LO
https://storage.googleapis.com/minikube/releases/latest/minikube_l
atest_amd64.deb
$ sudo dpkg -i minikube_latest_amd64.deb

$ sudo chmod 777 /var/run/docker.sock
$ sudo snap install kubectl --classic

kubectl 1.28.5 from Canonical✓ installed
$ minikube start
```

We have installed and Kubernetes.

```
ubuntu@ip-172-31-27-96:~$ minikube start
* minikube v1.32.0 on Ubuntu 20.04 (xen/amd64)
* Automatically selected the docker driver. Other choices: ssh, none
* Using Docker driver with root privileges
* Starting control plane node minikube in cluster minikube
* Pulling base image ...
* Downloading Kubernetes v1.28.3 preload ...
  > preloaded-images-k8s-v18-v1...: 403.35 MiB / 403.35 MiB 100.00% 15.66 M
  > gcr.io/k8s-minikube/kicbase...: 453.90 MiB / 453.90 MiB 100.00% 14.89 M
* Creating docker container (CPUs=2, Memory=2200MB) ...
* Preparing Kubernetes v1.28.3 on Docker 24.0.7 ...
  - Generating certificates and keys ...
  - Booting up control plane ...
  - Configuring RBAC rules ...
* Configuring bridge CNI (Container Networking Interface) ...
  - Using image gcr.io/k8s-minikube/storage-provisioner:v5
* Verifying Kubernetes components...
* Enabled addons: storage-provisioner, default-storageclass
* Done! kubectl is now configured to use "minikube" cluster and "default" namespace by default
```

Now developing the deployment config

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apache-deployment
  labels:
    app: apache
spec:
  replicas: 2
  selector:
    matchLabels:
      app: apache
  template:
    metadata:
      labels:
        app: apache
    spec:
      containers:
        - name: apache
          image: ubuntu/apache2
          ports:
            - containerPort: 80
```

Let's Deployment

```
$ kubectl create -f apache.yaml
deployment.apps/apache-deployment created
```

Check the deployment.

```
$ kubectl get deploy
NAME                 READY   UP-TO-DATE   AVAILABLE   AGE
apache-deployment    2/2     2            2           4m16s
```

```
$ kubectl expose deployment apache-deployment --type NodePort

$ kubectl get svc
NAME                                TYPE                CLUSTER-IP      EXTERNAL-IP
PORT(S)    AGE
apache-deployment    ClusterIP           10.104.65.184    <none>
80/TCP        3m22s
kubernetes           ClusterIP           10.96.0.1        <none>
443/TCP       26m
```

First Deployment has been completed.

```
$ git clone https://github.com/hshar/website.git
Cloning into 'website'...
remote: Enumerating objects: 8, done.
remote: Total 8 (delta 0), reused 0 (delta 0), pack-reused 8
Unpacking objects: 100% (8/8), 82.67 KiB | 7.52 MiB/s, done.

$ cd website
$ sudo nano Dockerfile
FROM ubuntu
RUN apt-get update
RUN apt-get install apache2 -y
ADD . /var/www/html
ENTRYPOINT apachectl -D FOREGROUND

$ sudo docker build . -t arkitcoin/k8scasestudy
Successfully built e14c9796d7e5
Successfully tagged arkitcoin/k8scasestudy:latest
```

```
ubuntu@ip-172-31-27-96:~/website$ docker images
REPOSITORY              TAG                IMAGE ID           CREATED            SIZE
arkitcoin/k8scasestudy  latest            e14c9796d7e5      About a minute ago 233MB
ubuntu                 latest            174c8c134b2a      4 weeks ago       77.9MB
ger.io/k8s-minikube/kicbase v0.0.42          dbc648475405      2 months ago      1.2GB
ubuntu@ip-172-31-27-96:~/website$ docker login
Login with your Docker ID to push and pull images from Docker Hub. If you don't have a Docker ID, head over to https://hub.docker.com to create one.
Username: arkitcoin
Password:
WARNING! Your password will be stored unencrypted in /home/ubuntu/.docker/config.json.
Configure a credential helper to remove this warning. See
https://docs.docker.com/engine/reference/commandline/login/#credentials-store

Login Succeeded
ubuntu@ip-172-31-27-96:~/website$ docker push arkitcoin/k8scasestudy
Using default tag: latest
The push refers to repository [docker.io/arkitcoin/k8scasestudy]
0563947fab86: Pushed
b19e0668c55e: Pushed
119db9557aa8: Pushed
a1360aae5271: Mounted from library/ubuntu
latest: digest: sha256:d8cd4d8034a5b6eb300bd68b57c1ad3aad9c925507c72557f2d6a62ad20646874 size: 1163
```

Create a custom deployment file

```
ubuntu@ip-172-31-27-96:~/website$ cat customdeploy.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: custom-deployment
  labels:
    app: custom
spec:
```

```
replicas: 2
selector:
  matchLabels:
    app: custom
template:
  metadata:
    labels:
      app: custom
  spec:
    containers:
      - name: custom
        image: arkitcoin/k8scasestudy
        ports:
          - containerPort: 80
```

```
$ kubectl create -f customdeploy.yaml
deployment.apps/custom-deployment created
```

```
$ kubectl expose deployment custom-deployment --type NodePort
service/custom-deployment exposed
```

If you want to check the website is working using Minikube just forward the port and access the website.

```
$ kubectl port-forward service/custom-deployment --address
0.0.0.0 :80
```

Deploy Ingress

```
$ kubectl get ingress
No resources found in default namespace.

$ minikube addons enable ingress
* ingress is an addon maintained by Kubernetes. For any
  concerns contact minikube on GitHub.
  You can view the list of minikube maintainers at:
  https://github.com/kubernetes/minikube/blob/master/OWNERS
  - Using image registry.k8s.io/ingress-nginx/kube-webhook-
    certgen:v20231011-8b53cabe0
  - Using image registry.k8s.io/ingress-
    nginx/controller:v1.9.4
  - Using image registry.k8s.io/ingress-nginx/kube-webhook-
    certgen:v20231011-8b53cabe0
* Verifying ingress addon...
* The 'ingress' addon is enabled
```

Lets define ingress

```

apiVersion: networking.k8s.io/v1
kind: Ingress
metadata:
  name: ingress
  annotations:
    nginx.ingress.kubernetes.io/rewrite-target: /
spec:
  ingressClassName: nginx
  rules:
  - http:
      paths:
      - path: /apache
        pathType: Prefix
        backend:
          service:
            name: apache-deployment
            port:
              number: 80
      - path: /custom
        pathType: Prefix
        backend:
          service:
            name: custom-deployment
            port:
              number: 80

```

```
$ kubectl create -f ingress.yaml
```

```
$ kubectl get ingress
```

NAME	CLASS	HOSTS	ADDRESS	PORTS	AGE
ingress	nginx	*	192.168.49.2	80	2m37s

```
$ kubectl get svc -A
```

NAMESPACE	NAME	TYPE
CLUSTER-IP	EXTERNAL-IP	PORT(S)
AGE		
default	apache-deployment	
ClusterIP	10.104.65.184	<none> 80/TCP
57m		
default	custom-deployment	NodePort
10.96.184.59	<none>	80:31900/TCP
27m		
default	kubernetes	
ClusterIP	10.96.0.1	<none> 443/TCP
80m		
ingress-nginx	ingress-nginx-controller	NodePort
10.108.87.183	<none>	80:30997/TCP, 443:31386/TCP
7m43s		

```
ingress-nginx    ingress-nginx-controller-admission
ClusterIP        10.99.101.192    <none>          443/TCP
7m43s
kube-system      kube-dns
ClusterIP        10.96.0.10       <none>
53/UDP,53/TCP,9153/TCP    80m
```

```
$ kubectl port-forward service/ingress-nginx-controller -n
ingress-nginx --address 0.0.0.0 :80
Forwarding from 0.0.0.0:32773 -> 80
```

Successfully Access the Website hosted on Kubernetes.