

## 3-Node Kubernetes Cluster Deployment

### Prerequisites:

- ✓ 3 virtual machines (VMs) are up and running in the same subnet / VPC. We are using EC2 instances on the AWS Cloud for this how-to guide. One of the VMs needs to have 2 vCPUs and 4 GB RAM, which means at least a t2.medium instance type. The rest of the VMs can be t2.micro instances.
- ✓ Ubuntu 22.04 LTS OS on all machines.
- ✓ Familiarity with the Kubernetes Components.
- ✓ On-Premises should have three VMs or Three Physical servers with communication enabled between them (Network)

### Part A - Controller and Worker Nodes

#### Configure Network Prerequisites

This step configures IPv4 forwarding and lets iptables see bridged traffic. The overlay module is needed by Docker to work with the overlay network driver. The br\_netfilter module is needed by Kubernetes to enable network filtering and NAT (Network Address Translation).

Refer to this [Kubernetes documentation](#) for more information.

```
cat <<EOF | sudo tee /etc/modules-load.d/k8s.conf
overlay
br_netfilter
EOF

sudo modprobe overlay
sudo modprobe br_netfilter

cat <<EOF | sudo tee /etc/sysctl.d/k8s.conf
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
EOF

sudo sysctl -system

lsmod | grep overlay
```

Confirm that you see an output similar to the following:

```
overlay                151552    0
```

```
lsmod | grep br_netfilter
```

Confirm that you see an output similar to the following:

```
br_netfilter           28672    0
bridge                 307200   1 br_netfilter
```

Verify that the `net.bridge.bridge-nf-call-iptables`, `net.bridge.bridge-nf-call-ip6tables`, and `net.ipv4.ip_forward` system variables are set to 1 in your `sysctl` config by running the following command:

```
sysctl net.bridge.bridge-nf-call-iptables net.bridge.bridge-nf-call-ip6tables net.ipv4.ip_forward
```

Output

```
net.bridge.bridge-nf-call-iptables = 1
net.bridge.bridge-nf-call-ip6tables = 1
net.ipv4.ip_forward = 1
```

Ensure to run these steps to configure the required network settings on all nodes.

## Configure Container Runtime

This step configures container runtime for all nodes on the cluster. Kubernetes 1.28 requires that you use a runtime that conforms with the Container Runtime Interface (CRI), which is an API for container run-times to work with `kubelet`. We will use `containerd` as the container runtime for this Kubernetes cluster.

Refer to [this Kubernetes documentation](#) for more information.

- i. Install Docker Engine (as it includes `containerd` runtime).

Refer to [this Docker documentation](#) for more information.

Update the `apt` package index, install the required packages, and get the official GPG keys for the Docker package repositories:

```
sudo apt-get update
sudo apt-get install ca-certificates curl gnupg
sudo install -m 0755 -d /etc/apt/keyrings
curl -fsSL https://download.docker.com/linux/ubuntu/gpg | sudo
gpg --dearmor -o /etc/apt/keyrings/docker.gpg
sudo chmod a+r /etc/apt/keyrings/docker.gpg
```

Add the Docker repository to the apt sources:

```
echo \
"deb [arch="$(dpkg --print-architecture)" signed-
by=/etc/apt/keyrings/docker.gpg]
https://download.docker.com/linux/ubuntu \
"$(. /etc/os-release && echo "$VERSION_CODENAME)" stable" | \
sudo tee /etc/apt/sources.list.d/docker.list > /dev/null
```

Update the apt package index, again:

```
sudo apt-get update
```

Install Docker (including containerd):

```
sudo apt-get install docker-ce docker-ce-cli containerd.io
docker-buildx-plugin docker-compose-plugin
```

Make daemon file for docker to avoid service errors later

```
sudo mkdir /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF

sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
```

Confirm that Docker Engine installation is successful by running the `hello-world` image:

```
sudo docker run hello-world
```

Switch-off the swap in all the nodes

```
sudo swapoff -a  
sudo sed -i '/ swap / s/^\(.*\)$/#\1/g' /etc/fstab
```

Configure the `containerd` runtime environment.

Refer to [this Kubernetes documentation](#) for more information.

Login as root:

```
sudo -i
```

Create a default `containerd` configuration file:

```
containerd config default > /etc/containerd/config.toml
```

Open `config.toml` in a text editor:

```
vi /etc/containerd/config.toml
```

Change the value of `SystemdCgroup` from `false` to `true` (it should be visible around line number 125 in `config.toml`):

```
SystemdCgroup = true
```

Restart `containerd`

```
systemctl restart containerd
```

Exit the `sudo` mode:

```
exit
```

Ensure to run these steps to configure the container runtime on **all** nodes.

Install kubeadm, kubelet, and kubect1

In this step, we install these packages on **all** nodes:

- kubeadm: The command to bootstrap the cluster
- kubelet: An agent that runs on all nodes in the cluster and does things like starting pods and containers
- kubect1: The command line utility to talk to the cluster

Refer to [this Kubernetes documentation](#) for more information.

Since we are using Ubuntu 22.04 LTS, we will use instructions for Debian-based distributions.

- i. Update the apt package index, install the required packages, and get the official GPG keys for the Kubernetes package repositories:

```
sudo apt-get update
sudo apt-get install -y apt-transport-https ca-certificates
curl
curl -fsSL
https://pkgs.k8s.io/core:/stable:/v1.28/deb/Release.key | sudo
gpg --dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
```

Add the Kubernetes repository to the apt sources:

```
echo 'deb [signed-by=/etc/apt/keyrings/kubernetes-apt-
keyring.gpg] https://pkgs.k8s.io/core:/stable:/v1.28/deb/ /' |
sudo tee /etc/apt/sources.list.d/kubernetes.list
```

Update the apt package index, again:

```
sudo apt-get update
```

Install kubelet, kubeadm, and kubect1:

```
sudo apt-get install -y kubelet kubeadm kubect1
```

Pin their version (avoids automatic updates):

```
sudo apt-mark hold kubelet kubeadm kubect1
```

Ensure to install kubeadm, kubelet, and kubectl on **all** nodes.

## Part B - Controller Node ONLY

Run these commands only on the VM designated as the controller (master) node.

(RUN AS ROOT) Initiate API server:

```
sudo -i
```

Initiate the API server. Remember to change the <ControllerVM-PrivateIP> with the actual private IP address of the controller VM:

```
kubeadm init --apiserver-advertise-address=<ControllerVM-PrivateIP> --  
pod-network-cidr=10.244.0.0/16
```

Exit the sudo mode

```
exit
```

(RUN AS NORMAL USER) Add a user for kube config:

```
mkdir -p $HOME/.kube
```

```
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
```

```
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

(RUN AS NORMAL USER) Deploy Weave network:

```
kubectl apply -f  
https://github.com/weaveworks/weave/releases/download/v2.8.1/weave-  
daemonset-k8s.yaml
```

(RUN AS ROOT) Create cluster join command:

```
sudo -i
```

```
kubeadm token create --print-join-command
```

## Part C - Worker Nodes ONLY

Copy the output of the cluster join command from the previous step and run on the VMs designated as the worker nodes.

(EXAMPLE COMMAND - DO NOT USE):

```
sudo kubeadm join 192.168.175.100:6443 --token jno9md.v9u1snltrwkv3vix  
--discovery-token-ca-cert-hash  
sha256:74b58d0e840e43a7051dcc4d7388b836dbd187c732fbfd3008051d10a14e271  
a
```

The Kubernetes cluster is now configured.

Now check the status of the nodes and cluster (RUN as NORMAL User)

```
kubect1 get nodes
```